

Python Pandas Tutorial: Series

A Series can be constructed with a plain list, dict, or scalar. Many operations on a Series have concise expression and are useful for powerful data analysis and munging.

Pandas is a powerful toolkit providing data analysis tools and structures for the Python programming language.

Among the most important artifacts provided by pandas is the Series. In this article, we introduce the Series class from a beginner's perspective. That means you do not need to know anything about pandas or data analysis to understand this tutorial.

What Is a Series?

A Series is similar to a list or an array in Python. It represents a series of values (numeric or otherwise) such as a column of data. Think of it as a Python list on steroids. It provides additional functionality, methods, and operators, which make it a more powerful version of a list.

To get started using a Series, you need to import the pandas toolkit into your Python program. Add the following line to the beginning of your program and you are good to go:

```
import pandas as pd
```

Create Series From List

Let's now learn how to create a Series. Since a Series is similar to a list, let's use a list to create a Series.

```
ser = pd.Series([1, 3, 5, 7])
```

```
print ser
```

```
0    1
```

```
1    3
```

```
2    5
```

```
3    7
```

```
dtype: int64
```

Notice that when the Series is printed, two columns of numbers are printed. The first

column is called the index. It normally starts from 0 and runs all the way to $(N-1)$ where N is the size of the list.

The second column is the actual data that we specified.

4. Use a Dict to Initialize Series

It is also possible to use a Python dictionary (called a dict) to initialize a Series. In this case, the Series takes its index from the keys of the dict as shown by this example.

```
prices = {'apple': 4.99,  
          'banana': 1.99,  
          'orange': 3.99,  
          'grapes': 0.99}
```

```
ser = pd.Series(prices)
```

```
print ser
```

```
apple    4.99
```

```
banana    1.99
```

```
grapes    0.99
```

```
orange    3.99
```

```
dtype: float64
```

5. Initialize Series from Scalar

You can also use a scalar to initialize a Series. In this case, all elements of the Series are initialized to the same value. When used with a scalar for initialization, an index array can be specified. In this case, the size of the Series is the same as the size of the index array.

In the following example, we use the `range()` function to specify the index (and thus the size of the Series).

```
ser = pd.Series(2, index=range(0, 5))
```

```
print ser
```

```
0    2
```

```
1    2
```

```
2    2
```

```
3    2
```

```
4    2
```

```
dtype: int64
```

6. Some Other Ways of Creating a Series

Here are some additional ways of initializing a Series.

A Series of Odd Numbers

We use the `range()` function to create a Series of odd numbers.

```
print pd.Series(range(1, 10, 2))
```

```
0    1
```

```
1    3
```

```
2    5
```

```
3    7
```

```
4    9
```

```
dtype: int64
```

With an Alphabetic Index

```
print pd.Series(range(1, 15, 3), index=[x for x in 'abcde'])
```

```
a     1
```

```
b     4
```

```
c     7
```

```
d    10
```

```
e    13
```

```
dtype: int64
```

A Series With Random Numbers

Use a random range to initialize a Series:

```
print pd.Series(random.sample(xrange(100), 6))
```

```
0    61
```

```
1    81
```

```
2    11
```

```
3    78
```

```
4    29
```

```
5    92
```

```
dtype: int64
```

Combining Lists

If you have the data in one list and the index in another list, there are a couple of ways to create a Series from them.

Using a Dict

```
x = dict(zip([x for x in 'abcdefg'], xrange(1, 8)))

print x

y = pd.Series(x)

print y

{'a': 1, 'c': 3, 'b': 2, 'e': 5, 'd': 4, 'g': 7, 'f': 6}

a    1
b    2
c    3
d    4
e    5
f    6
g    7

dtype: int64
```

Specifying an Index

Skip defining a dict and create the Series directly.

```
print pd.Series(xrange(1,8), index=[x for x in 'abcdefg'])

a    1
b    2
c    3
d    4
e    5
f    6
g    7

dtype: int64
```

7. Naming a Series

You can also assign a name to the Series. When used in the context of a DataFrame (to be covered next), this name serves as the column name.

```
a = [1, 3, 5, 7]
```

```
print pd.Series(a, name='joe')
```

```
0    1
```

```
1    3
```

```
2    5
```

```
3    7
```

```
Name: joe, dtype: int64
```

8. Comparing List with Series

A Series is like a list. The usual array indexing works as expected, as does the array slicing. Notice that the slice operator returns a Series itself.

```
ser = pd.Series(random.sample(xrange(100), 10))
```

```
print ser
```

```
print
```

```
print '4th element: ', ser[4]
```

```
print 'Slice: ', ser[3:8]
```

```
0    79
```

```
1     4
```

```
2    71
```

```
3    20
```

```
4    19
```

```
5    24
```

```
6    82
```

```
7    74
```

```
8    17
```

```
9    48
```

```
dtype: int64
```

```
4th element:  19
```

```
Slice:  3    20
```

```
4    19
```

```
5    24
```

```
6    82
```

```
7    74
```

```
dtype: int64
```

Using a Predicate for Extraction

While list elements can be extracted using indexing and slicing, Series elements can also be extracted using a predicate function (a function returning `True` or `False`) as shown in this example.

```
x = random.sample(range(100), 10)

print x

y = pd.Series(x)

print y[y > 30]

[22, 5, 16, 56, 38, 48, 64, 41, 71, 63]

3    56
4    38
5    48
6    64
7    41
8    71
9    63

dtype: int64
```

Indexing Using a List

In addition to extraction using a predicate function, items can also be retrieved from a Series using a list as the index.

```
x = random.sample(xrange(100), 10)

print x

y = pd.Series(x)

print y[[2, 0, 1, 2]]

[29, 1, 27, 54, 2, 90, 25, 96, 45, 34]

2    27
0    29
1     1
2    27

dtype: int64
```

Executing a Function on a Series

Unlike a list, a function accepting a scalar argument can be invoked with a Series. The result will be another Series with the function applied to each element of the Series. This allows more flexible and concise ways in which operations can be combined — just what is needed in a data analysis toolkit!

```
def joe(x): return x + 10

x = random.sample(xrange(100), 10)

print 'Data => ', x, '\n'

y = pd.Series(x)

print 'Applying pow => \n', pow(y, 2), '\n'

print 'Applying joe => \n', joe(y)

Data => [96, 63, 79, 11, 49, 41, 12, 26, 20, 62]

Applying pow =>

0    9216
1    3969
2    6241
3     121
4    2401
5    1681
6     144
7     676
8     400
9    3844

dtype: int64

Applying joe =>

0    106
1     73
2     89
3     21
4     59
5     51
6     22
7     36
8     30
```

```
dtype: int64
```

9. Comparing a Dict With a Series

A Series also behaves like a Python dictionary.

Indexing With Label

For example, you can extract items using the index label.

```
x = pd.Series(xrange(1,8), index=[x for x in 'abcdefg'])
```

```
print x, '\n'
```

```
print 'element "d" => ', x['d']
```

```
a    1
```

```
b    2
```

```
c    3
```

```
d    4
```

```
e    5
```

```
f    6
```

```
g    7
```

```
dtype: int64
```

```
element "d" => 4
```

Checking for Membership

Use the `in` operator to check whether a Series contains a specific label.

```
x = pd.Series(xrange(1,8), index=[x for x in 'abcdefg'])
```

```
print x, '\n'
```

```
print 'is "j" in x?', 'j' in x
```

```
print 'is "d" in x?', 'd' in x
```

```
a    1
```

```
b    2
```

```
c    3
```

```
d    4
```

```
e    5
```

```
f    6
```



```
g      7
```

```
dtype: int64
```

```
is "j" in x? False
```

```
is "d" in x? True
```

Summary

This article was a gentle introduction to pandas Series. A Series can be constructed using a plain list, a dict, or a scalar. Many operations on a Series allow concise expression and are useful for powerful data analysis and munging.

Python Pandas Tutorial: Series Methods

Learn commonly used methods to deal with a Series object, including methods to retrieve general information about a Series, modifying a Series, selection, and sorting.

The Series is one of the most common Pandas data structures. It is similar to a Python list and is used to represent a column of data. After looking into the basics of [creating and initializing](#) a pandas Series object, we now delve into some common usage patterns and methods.

Series Information

After a Series is created, it is most important to look into various details of its structure. These include the size of the series, whether there are NaNs in it, etc. Here are some commonly used methods which help clarify the situation.

Size of the Series

There are several methods to determine how big the Series is.

The first is the attribute `shape`, which returns a tuple.

```
a = pd.Series(random.sample(xrange(100), 6))  
  
print a.shape  
  
(6,)
```

We also have the `count()` method which returns the size of the Series as an integer.

```
print a.count()  
  
6
```

However, note that `count()` only reports the number of non-NaN elements, while `shape` reports both.

Another attribute for getting the count of elements is `size`. It reports the count as an integer and includes NaN elements if any.

```
a = pd.Series(random.sample(xrange(100), 6))  
  
print 'count of a =>', a.count(), '\n'
```

```

b = a.append(pd.Series(np.nan, index=list('abcd')), ignore_index=True)

print 'b => ', b, '\n'

print 'count of b =>', b.count(), '\n'

print 'shape of b =>', b.shape, '\n'

print 'size of b =>', b.size

count of a => 6

b =>  0    76.0

1    92.0

2    75.0

3    60.0

4    42.0

5    44.0

6     NaN

7     NaN

8     NaN

9     NaN

dtype: float64

count of b => 6

shape of b => (10,)

size of b => 10

```

Series Details

Get some detailed stats on the Series using `describe()` . This method returns a Series object with the index (or labels) as shown.

```

x = pd.Series(random.sample(xrange(100), 6))

x.describe()

count      6.000000
mean       60.500000
std        30.742479
min        20.000000
25%        44.000000
50%        53.500000
75%        86.250000

```

```
max      98.000000
```

```
dtype: float64
```

Head and Tail

Show the first 5 or last 5 rows of the Series using `head()` or `tail()` .

```
x = pd.Series(random.sample(xrange(100), 10))
```

```
print x, '\n'
```

```
print x.head(), '\n'
```

```
print x.tail(), '\n'
```

```
0      24
```

```
1      39
```

```
2      56
```

```
3      77
```

```
4      81
```

```
5      26
```

```
6       8
```

```
7      87
```

```
8      34
```

```
9      68
```

```
dtype: int64
```

```
0      24
```

```
1      39
```

```
2      56
```

```
3      77
```

```
4      81
```

```
dtype: int64
```

```
5      26
```

```
6       8
```

```
7      87
```

```
8      34
```

```
9      68
```

```
dtype: int64
```

Add Elements to Series

Adding elements to a Series is accomplished by using `append()`. The argument must be a single Series object, or a list (or tuple) of Series objects.

```
x = pd.Series(random.sample(xrange(100), 6))

print x, '\n'

print 'appended =>\n', x.append([pd.Series(2), pd.Series([3, 4, 5])])

0    62
1    29
2    20
3    69
4    53
5    22

dtype: int64

appended =>

0    62
1    29
2    20
3    69
4    53
5    22
0     2
0     3
1     4
2     5

dtype: int64
```

You might notice the oddball labels after appending. Each Series is appended with a default index starting from 0, regardless of whether this creates duplicate labels. One way to fix this is to specify `ignore_index=True` to ensure re-labeling.

```
print 'appended =>\n', x.append([pd.Series(2), pd.Series([3, 4, 5])],
ignore_index=True)

appended =>

0    62
1    29
```

```
2    20
3    69
4    53
5    22
6     2
7     3
8     4
9     5
```

```
dtype: int64
```

What if you don't want to re-label but ensure that `append()` succeeds only if the labels are unique? Keep your precious labels intact and unique by specifying `verify_integrity=True`.

```
print 'appended =>\n', x.append([pd.Series(2), pd.Series([3, 4, 5])],
verify_integrity=True)
```

```
ValueError: Indexes have overlapping values: [0, 1, 2]
```

Delete Elements

You can delete elements from a Series using the following methods.

By Label

Use `drop()` and specify a single label or a list of labels to drop.

```
x = pd.Series(random.sample(xrange(100), 6), index=list('ABCDEF'))
```

```
print x, '\n'
```

```
print 'drop one =>\n', x.drop('C'), '\n'
```

```
print 'drop many =>\n', x.drop(['C', 'D'])
```

```
A    67
B    18
C     1
D    54
E    38
F     3
```

```
dtype: int64
```

```
drop one =>
```

```
A    67
```

B 18

D 54

E 38

F 3

dtype: int64

drop many =>

A 67

B 18

E 38

F 3

dtype: int64

Duplicate Elements

Get rid of duplicate elements by invoking `drop_duplicates()` .

```
x = pd.Series([1, 2, 2, 4, 5, 7, 3, 4])
```

```
print x, '\n'
```

```
print 'drop duplicates =>\n', x.drop_duplicates(), '\n'
```

0 1

1 2

2 2

3 4

4 5

5 7

6 3

7 4

dtype: int64

drop duplicates =>

0 1

1 2

3 4

4 5

5 7

6 3

```
dtype: int64
```

By default, the method retains the first repeated value. Get rid of all duplicates (including the first) by specifying `keep=False`.

```
drop all duplicates =>
```

```
0    1
```

```
4    5
```

```
5    7
```

```
6    3
```

```
dtype: int64
```

NaN Elements

Use the `dropna()` to drop elements without a value (NaN).

```
x = pd.Series([1, 2, 3, 4, np.nan, 5, 6])
```

```
print x, '\n'
```

```
print 'drop na =>\n', x.dropna()
```

```
0    1.0
```

```
1    2.0
```

```
2    3.0
```

```
3    4.0
```

```
4    NaN
```

```
5    5.0
```

```
6    6.0
```

```
dtype: float64
```

```
drop na =>
```

```
0    1.0
```

```
1    2.0
```

```
2    3.0
```

```
3    4.0
```

```
5    5.0
```

```
6    6.0
```

```
dtype: float64
```

Replace NaN Elements

When you want to replace NaN elements in a Series, use `fillna()` .

```
x = pd.Series([1, 2, 3, 4, np.nan, 5, 6])
```

```
print x, '\n'
```

```
print 'fillna w/0 =>\n', x.fillna(0)
```

```
0    1.0
```

```
1    2.0
```

```
2    3.0
```

```
3    4.0
```

```
4    NaN
```

```
5    5.0
```

```
6    6.0
```

```
dtype: float64
```

```
fillna w/0 =>
```

```
0    1.0
```

```
1    2.0
```

```
2    3.0
```

```
3    4.0
```

```
4    0.0
```

```
5    5.0
```

```
6    6.0
```

```
dtype: float64
```

Select Elements

Select elements from a Series based on various conditions as follows.

In a Range

Use the `between()` method, which returns a Series of boolean values indicating whether the element lies within the range.

```
a = pd.Series(random.sample(xrange(100), 10))
```

```
print a
```

```
print a.between(30, 50)
```

```
0    85
```

```
1    42
```

```
2    63
3    69
4    81
5    45
6    50
7    72
8    66
9    34
```

```
dtype: int64
```

```
0    False
1     True
2    False
3    False
4    False
5     True
6     True
7    False
8    False
9     True
```

```
dtype: bool
```

You can use this the returned boolean Series as a predicate into the original Series.

```
print a[a.between(30, 50)]
```

```
1    42
5    45
6    50
9    34
```

```
dtype: int64
```

Using a Function

Select elements using a predicate function as the argument to `select()` .

```
x = pd.Series(random.sample(xrange(100), 6))
```

```
print x, '\n'
```

```

print 'select func =>\n', x.select(lambda a: x.iloc[a] > 20)

0      83
1      96
2      29
3      15
4      28
5      12

dtype: int64

select func =>

0      83
1      96
2      29
4      28

dtype: int64

```

By List of Labels

Use `filter(items=[..])` with the labels to be selected in a list.

```

x = pd.Series([1, 2, 3, 4, np.nan, 5, 6])

print x, '\n'

print 'filtered =>\n', x.filter(items=[1, 2, 6])

0      1.0
1      2.0
2      3.0
3      4.0
4      NaN
5      5.0
6      6.0

dtype: float64

filtered =>

1      2.0
2      3.0
6      6.0

dtype: float64

```

Regex Match on Label

Select labels to filter using a regular expression match with `filter(regex='...')`.

```
x = pd.Series({'apple': 1.99,
               'orange': 2.49,
               'banana': 0.99,
               'grapes': 1.49,
               'melon': 3.99})

print x, '\n'

print 'regex filter =>\n', x.filter(regex='a
)

apple      1.99
banana     0.99
grapes     1.49
melon      3.99
orange     2.49
dtype: float64

regex filter =>

banana     0.99
dtype: float64
```

Substring Match on Label

Use the `filter(like='...')` version to perform a substring match on the labels to be selected.

```
print 'like filter =>\n', x.filter(like='an')

like filter =>

banana     0.99
orange     2.49
dtype: float64
```

Sorting

Ah! Sorting. The all important functionality when playing with data.

Here is how you can sort a Series by labels or by value.

By Index (or Labels)

Use `sort_index()` .

```
x = pd.Series(random.sample(xrange(100), 6), index=random.sample(map(chr,
xrange(ord('a'), ord('z'))), 6))
```

```
print x, '\n'
```

```
print 'sort by index: =>\n', x.sort_index(), '\n'
```

```
p    37
```

```
e    44
```

```
b    93
```

```
l    75
```

```
n     4
```

```
s    83
```

```
dtype: int64
```

```
sort by index: =>
```

```
b    93
```

```
e    44
```

```
l    75
```

```
n     4
```

```
p    37
```

```
s    83
```

```
dtype: int64
```

By Values

Use `sort_values()` to sort by the values.

```
print 'sort by value: =>\n', x.sort_values()
```

```
sort by value: =>
```

```
n     4
```

```
p    37
```

```
e    44
```

```
l    75
```

```
s    83
```

```
b    93
```

```
dtype: int64
```

Summary

This article covers some commonly used methods to deal with a Series object, including methods to retrieve general information about a Series, modifying a Series, selection, and sorting.

Python Pandas Tutorial: DataFrame Basics

The most commonly used data structures in pandas are DataFrames, so it's important to know at least the basics of working with them.

The DataFrame is the most commonly used data structures in pandas. As such, it is very important to learn various specifics about working with the DataFrame. After of creating a DataFrame, let's now delve into some methods for working with it.

Getting Started

Import these libraries: `pandas` , `matplotlib` for plotting, `numpy` .

```
import pandas as pd

import matplotlib.pyplot as plt

import numpy as np

import random
```

If you are working with a Jupyter (or iPython) notebook and want to show graphs inline, use this definition.

```
%matplotlib inline
```

Let's now load some CSV data into our DataFrame for working with it. The data we have loaded is the [World Happiness Report 2016](#).

```
x = pd.read_csv('2016.csv')
```

DataFrame Details

Index

The attribute `index` shows the row index labels.

```
x = pd.read_csv('2016.csv')

print x.index
```

```
RangeIndex(start=0, stop=157, step=1)
```

The index is a `RangeIndex` if the labels are contiguous integers.

Columns

Get the columns using the attribute `columns` .

```
print x.columns  
  
Index([u'Country', u'Region', u'Happiness Rank', u'Happiness Score',  
       u'Lower Confidence Interval', u'Upper Confidence Interval',  
       u'Economy (GDP per Capita)', u'Family', u'Health (Life Expectancy)',  
       u'Freedom', u'Trust (Government Corruption)', u'Generosity',  
       u'Dystopia Residual'],  
      dtype='object')
```

Values

The raw values array can be extracted using `values` .

```
print x.values  
  
[['Denmark' 'Western Europe' 1 ..., 0.444530000000000004 0.36171 2.73939]  
 ['Switzerland' 'Western Europe' 2 ..., 0.41203 0.28083 2.69463]  
 ['Iceland' 'Western Europe' 3 ..., 0.14975 0.476780000000000004 2.83137]  
 ...,  
 ['Togo' 'Sub-Saharan Africa' 155 ..., 0.11587 0.17517 2.1354]  
 ['Syria' 'Middle East and Northern Africa' 156 ..., 0.17232999999999998  
  0.48396999999999996 0.81789]  
 ['Burundi' 'Sub-Saharan Africa' 157 ..., 0.09419 0.2029 2.1040400000000004]]
```

Shape

Get a tuple of the number of rows and columns of the DataFrame using the `shape` attribute.

```
x.shape  
  
(157, 13)
```

Size

Use the `count()` method to retrieve a count of (non-NaN) elements in each column. This method ignores any NaN elements in the column.

```
print x.count()  
  
Country          157  
Region           157
```


Happiness Rank	157
Happiness Score	157
Lower Confidence Interval	157
Upper Confidence Interval	157
Economy (GDP per Capita)	157
Family	157
Health (Life Expectancy)	157
Freedom	157
Trust (Government Corruption)	157
Generosity	157
Dystopia Residual	157

dtype: int64

And the `size` attribute returns the total number of elements (including NaNs) in the DataFrame. This means the value (nrows * ncols).

```
print x.size
```

2041

Statistics

Get detailed statistics of the DataFrame using the method `describe()`. Returns various details such as mean, min, max, etc. for each column.

```
print x.describe()
```

	Happiness Rank	Happiness Score	Lower Confidence Interval	\
count	157.000000	157.000000	157.000000	
mean	78.980892	5.382185	5.282395	
std	45.466030	1.141674	1.148043	
min	1.000000	2.905000	2.732000	
25%	40.000000	4.404000	4.327000	
50%	79.000000	5.314000	5.237000	
75%	118.000000	6.269000	6.154000	
max	157.000000	7.526000	7.460000	

	Upper Confidence Interval	Economy (GDP per Capita)	Family	\
count	157.000000	157.000000	157.000000	
mean	5.481975	0.953880	0.793621	

```
std                                1.136493                0.412595        0.266706
...

```

Head and Tail

The `head()` method retrieves the first five rows from the DataFrame.

```
x = pd.read_csv('big-data/Salaries.csv')
print x.head()
```

	yearID	teamID	lgID	playerID	salary
0	1985	ATL	NL	barkele01	870000
1	1985	ATL	NL	bedrost01	550000
2	1985	ATL	NL	benedbr01	545000
3	1985	ATL	NL	campri01	633333
4	1985	ATL	NL	ceronri01	625000

And the tail method retrieves the last five rows.

```
print x.tail()
```

	yearID	teamID	lgID	playerID	salary
26423	2016	WSN	NL	strasst01	10400000
26424	2016	WSN	NL	taylomi02	524000
26425	2016	WSN	NL	treinbl01	524900
26426	2016	WSN	NL	werthja01	21733615
26427	2016	WSN	NL	zimmer01	14000000

The cumulative methods return a DataFrame with the appropriate cumulative function applied to the rows. Some of the operations are not valid for non-numeric columns.

Cumulative Sum

`cumsum()` (cumulative sum): Value of each row is replaced by the sum of all prior rows including this row. String value rows use concatenation as shown below.

```
y = pd.DataFrame({'one': pd.Series(range(5)),
                  'two': pd.Series(range(5, 10)),
                  'three': pd.Series(list('abcde'))})

print 'head =>\n', y.head(), '\n'

print 'cumsum =>\n', y.cumsum(), '\n'

head =>
```

	one	three	two
0	0	a	5
1	1	b	6
2	2	c	7
3	3	d	8
4	4	e	9

cumsum =>

	one	three	two
0	0	a	5
1	1	ab	11
2	3	abc	18
3	6	abcd	26
4	10	abcde	35

Cumulative Product

cumprod() (**cumulative product**): Row value is replaced by product of all prior rows. This method is not applicable to non-numeric rows. If there are non-numeric rows in the DataFrame, you will need to extract a subset of the DataFrame as shown.

```
print 'cumprod =>\n', y[['one', 'two']].cumprod(), '\n'
```

cumprod =>

	one	two
0	0	5
1	0	30
2	0	210
3	0	1680
4	0	15120

Cumulative Maximum

cummax() (**cumulative max**): Value of the row is replaced by the maximum value of all prior rows till now. In the example below, for demonstrating this method, we use this method on reversed rows of the original DataFrame.

```
print 'rev =>\n', y.iloc[::-1], '\n',
```

```
print 'cummax =>\n', y.iloc[::-1].cummax(), '\n'
```

rev =>

	one	three	two
4	4	e	9
3	3	d	8
2	2	c	7
1	1	b	6
0	0	a	5

cummax =>

	one	three	two
4	4	e	9
3	4	e	9
2	4	e	9
1	4	e	9
0	4	e	9

Cumulative Minimum

cummin() : Similar to cummax, except computes the minimum of values till this row.

```
print 'cummin =>\n', y.cummin(), '\n'
```

cummin =>

	one	three	two
0	0	a	5
1	0	a	5
2	0	a	5
3	0	a	5
4	0	a	5

Index of Min and Max Values

Use the methods **idxmin()** and **idxmax()** to obtain the index label of the rows containing minimum and maximum values. Applicable only to numeric columns, so non-numeric columns need to be filtered out.

```
y = pd.DataFrame({'one': pd.Series(random.sample(xrange(100), 5),
index=list('abcde')),

                  'two': pd.Series(random.sample(xrange(100), 5),
index=list('abcde')),

                  'three': pd.Series(list('ABCDE'), index=list('abcde'))})
```

```

print y, '\n'

print 'idxmax =>\n', y[['one', 'two']].idxmax(), '\n'

print 'idxmin =>\n', y[['one', 'two']].idxmin(), '\n'

    one three  two
a    48      A   25
b    38      B   13
c    62      C   91
d    79      D   32
e     2      E   42

idxmax =>

one      d

two      c

dtype: object

idxmin =>

one      e

two      b

dtype: object

```

Value Counts

The method `value_counts()` returns the number of times each value is repeated in the column. Note: this is not a DataFrame method; rather it is applicable on a column (which is a Series object).

```

x = pd.read_csv('big-data/Salaries.csv')

print 'top 10 =>\n', x.head(10), '\n'

print 'value_counts =>\n', x['yearID'].value_counts().head(10)

top 10 =>

```

	yearID	teamID	lgID	playerID	salary
0	1985	ATL	NL	barkeler01	870000
1	1985	ATL	NL	bedrost01	550000
2	1985	ATL	NL	benedbr01	545000
3	1985	ATL	NL	campri01	633333
4	1985	ATL	NL	ceronri01	625000
5	1985	ATL	NL	chambch01	800000

```
6    1985    ATL    NL    dedmoje01    150000
7    1985    ATL    NL    forstte01    483333
8    1985    ATL    NL    garbege01    772000
9    1985    ATL    NL    harpete01    250000
```

```
value_counts =>
```

```
1999    1006
1998     998
1995     986
1996     931
1997     925
1993     923
1994     884
1990     867
2001     860
2008     856
```

```
Name: yearID, dtype: int64
```

Summary

We covered a few aspects of the DataFrame in this article. Ways of learning various details of the DataFrame including size, shape, statistics, etc. were presented.